



## ML and DL

635 | Vishal Baraiya |  
23010101014

## Lab - 6

### Decision Tree Classifier

```
In [1]: from sklearn.datasets import load_iris
```

### Import iris data ser using sklearn

```
In [2]: data = load_iris()
```

```
In [3]: type(data)
```

```
Out[3]: sklearn.utils._bunch.Bunch
```

### Importing the libraries

```
In [4]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
```

### Divide the data into input and output

```
In [5]: x = data["data"]
```

```
In [6]: y = data["target"]  
y
```

# Splitting the dataset into the Training set and Test set

```
In [7]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=42)
```

## Fitting DecisionTreeClassifier on dataset

```
In [8]: model = DecisionTreeClassifier(random_state=42)
```

```
In [9]: model.fit(x_train, y_train)
```

Out[9]: `DecisionTreeClassifier`

```
DecisionTreeClassifier(random_state=42)
```

# Display Decision Tree

In [10]: %whos

Variable	Type	Data/Info
DecisionTreeClassifier	ABCMeta	<class 'sklearn.tree._cla<...>.De
cisionTreeClassifier'>		
data	Bunch	{'data': array([[5.1, 3.5<...> 's
klearn.datasets.data'}		
load_iris	function	<function load_iris at 0x00000185
8702A200>		
model	DecisionTreeClassifier	DecisionTreeClassifier(random_st
te=42)		
np	module	<module 'numpy' from 'C:\<...>ges
\\"numpy\\__init__.py'>		
pd	module	<module 'pandas' from 'C:<...>es
\\"pandas\\__init__.py'>		
plt	module	<module 'matplotlib.pyplot.py'>\\"m
atplotlib\\pyplot.py'>		
sns	module	<module 'seaborn' from 'C<...>s
\\"seaborn\\__init__.py'>		
train_test_split	function	<function train_test_split at 0x0
00001858A636F20>		
tree	module	<module 'sklearn.tree' fr<...>ear
n\\tree\\__init__.py'>		
x	ndarray	150x4: 600 elems, type `float64` ,
4800 bytes		
x_test	ndarray	45x4: 180 elems, type `float64` ,
1440 bytes		
x_train	ndarray	105x4: 420 elems, type `float64` ,
3360 bytes		
y	ndarray	150: 150 elems, type `int32` , 600
bytes		
y_test	ndarray	45: 45 elems, type `int32` , 180 b
ytes		
y_train	ndarray	105: 105 elems, type `int32` , 420
bytes		

In [11]: %matplotlib qt

```
In [12]: plt.figure(figsize=(20, 10))
tree.plot_tree(
    model,
    filled=True,
    feature_names=data.feature_names,
    class_names=data.target_names,
    rounded=True
)
plt.show()
```

## Predict the x\_test

In [13]: y\_pred = model.predict(x\_test)  
y\_pred

```
Out[13]: array([1, 0, 2, 1, 1, 0, 1, 2, 1, 1, 2, 0, 0, 0, 0, 1, 2, 1, 1, 2, 0, 2, 2, 2, 2, 2, 0, 0, 0, 1, 0, 0, 2, 1, 0, 0, 0, 2, 1, 1, 0, 0])
```

```
In [14]: from sklearn.metrics import accuracy_score  
  
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 1.0

## Import diabetes.csv dataset

```
In [15]: df = pd.read_csv('diabetes.csv')  
df
```

```
Out[15]:    Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  DiabetesPedigreeFunction  Age  
0             6        148              72                35       0   33.6  
1             1         85              66                29       0   26.6  
2             8        183              64                0       0   23.3  
3             1         89              66                23     94   28.1  
4             0        137              40                35    168   43.1  
...           ...       ...            ...            ...       ...  ...  ...  
763            10        101              76                48    180   32.9  
764            2        122              70                27       0   36.8  
765            5        121              72                23    112   26.2  
766            1        126              60                0       0   30.1  
767            1         93              70                31       0   30.4
```

768 rows × 9 columns



## Check the distribution of the target

```
In [16]: print(df['Outcome'].value_counts())
```

```
Outcome
Non Diabetic    500
Diabetic        268
Name: count, dtype: int64
```

## Replace zeros with NaN for specific columns

```
In [17]: columns_to_replace = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']
df[columns_to_replace] = df[columns_to_replace].replace(0,np.nan)
```

## Check for missing values

```
In [18]: print(df.isnull().sum())
```

```
Pregnancies      0
Glucose          5
BloodPressure    35
SkinThickness   227
Insulin          374
BMI              11
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
```

## Fill missing values with median

```
In [19]: df.fillna(df.median(numeric_only=True), inplace=True)
```

```
In [20]: print(df.isnull().sum())
```

```
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness   0
Insulin          0
BMI              0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
```

## Visualize Distributions

```
In [21]: df.hist(figsize=(12, 10), bins=20)
plt.tight_layout()
```

```
plt.show()
```

```
In [22]: sns.countplot(x='Outcome', data=df)
plt.title('Distribution of Outcome')
plt.show()
```

## Convert Targer data into interger code

```
In [23]: df['Outcome'] = df['Outcome'].map({'Non Diabetic': 0, 'Diabetic': 1})
```

```
In [24]: df.head()
```

```
Out[24]:    Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  DiabetesPedigreeFunc
0             6        148.0          72.0            35.0     125.0   33.6
1             1         85.0          66.0            29.0     125.0   26.6
2             8        183.0          64.0            29.0     125.0   23.3
3             1         89.0          66.0            23.0      94.0   28.1
4             0        137.0          40.0            35.0     168.0   43.1
```

## Divide the data into input and output

```
In [25]: X = df.drop('Outcome', axis=1)
X
```

Out[25]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148.0	72.0	35.0	125.0	33.6	
1	1	85.0	66.0	29.0	125.0	26.6	
2	8	183.0	64.0	29.0	125.0	23.3	
3	1	89.0	66.0	23.0	94.0	28.1	
4	0	137.0	40.0	35.0	168.0	43.1	
...	...	...	...	...	...	...	...
763	10	101.0	76.0	48.0	180.0	32.9	
764	2	122.0	70.0	27.0	125.0	36.8	
765	5	121.0	72.0	23.0	112.0	26.2	
766	1	126.0	60.0	29.0	125.0	30.1	
767	1	93.0	70.0	31.0	125.0	30.4	

768 rows × 8 columns



In [26]: `Y = df['Outcome']`  
Y

Out[26]:

0	1
1	0
2	1
3	0
4	1
..	..
763	0
764	0
765	0
766	1
767	0

Name: Outcome, Length: 768, dtype: int64

## Splitting the dataset into the Training set and Test set

In [27]: `X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3, random_state=42)`

## Create Model

In [28]: `dt_model = DecisionTreeClassifier(random_state=42)`

# Fitting DecisionTreeClassifier on dataset

```
In [29]: dt_model.fit(X_train, Y_train)
```

Out[29]:

```
▼      DecisionTreeClassifier ⓘ ⓘ
DecisionTreeClassifier(random_state=42)
```

## Display Decision Tree

```
In [30]: plt.figure(figsize=(25, 15))
tree.plot_tree(
    dt_model,
    filled=True,
    feature_names=X.columns,
    class_names=['No Diabetes', 'Diabetes'],
    rounded=True
)
plt.show()
```

## Predict the x\_test

```
In [31]: Y_pred = dt_model.predict(X_test)
Y_pred
```

```
Out[31]: array([0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0,
0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1,
0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1,
0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1,
0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1,
1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
```

## Display Training Accuracy

```
In [32]: train_accuracy = accuracy_score(Y_train, dt_model.predict(X_train))
print("Training Accuracy:", train_accuracy)
```

Training Accuracy: 1.0

## Display Test Accuracy

```
In [33]: test_accuracy = accuracy_score(Y_test, Y_pred)
print("Test Accuracy:", test_accuracy)
```

```
Test Accuracy: 0.696969696969697
```

## Confusion Matrix

```
In [34]: from sklearn.metrics import confusion_matrix
```

```
In [35]: cm = confusion_matrix(Y_test, Y_pred)
print("Confusion Matrix:")
print(cm)
```

```
Confusion Matrix:
```

```
[[113  38]
 [ 32  48]]
```

```
In [36]: plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Diabetes', 'Diabetes'],
            yticklabels=['No Diabetes', 'Diabetes'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

## Save Model State

```
In [37]: import pickle
```

```
In [38]: with open('model.pkl', 'wb') as file:
    pickle.dump(model, file)
```

## Load Model from Disk

```
In [39]: with open('model.pkl', 'rb') as file:
    loaded_model = pickle.load(file)
```

```
In [40]: y_predict = loaded_model.predict(x_test)
```

```
In [41]: from sklearn.metrics import classification_report
```

```
In [42]: print(classification_report(y_test, y_predict))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

In [ ]: